INTEL UNNATI 5 PAGE REPORT

BY CHINTHAGUNTA VAMSHI KRISHNA

INSTITUTE OF AERONAUTICAL ENGINEERING (IARE)

BRANCH- CSE (AIML) – 22951A66G2

Fine-Tuning a Large Language Model to Create a Custom Chatbot.

1. Introduction

Problem Statement

The goal of this project is to fine-tune a large language model (LLM) to create a custom chatbot using readily available hardware, specifically 4th Generation Intel® Xeon® Scalable processors. Participants will use a systematic methodology to generate a domain-specific dataset and optimize the fine-tuning process with Intel® Extension for Transformers' Neural Chat.

Objectives

1. Train and fine-tune a custom chatbot.
2. Utilize the Intel Developer Cloud (IDC) for development and deployment.
3. Implement fine-tuning using the Alpaca Dataset and Llama 2 model.

2. Technical Approach

Dataset

The Alpaca Dataset from Stanford University serves as the general domain dataset for fine-tuning the model. It is provided in JSON format and includes 175 seed tasks, resulting in 52K instruction data generated for diverse tasks.

Model

Llama 2 is a family of pre-trained and fine-tuned large language models developed by Meta, ranging from 7B to 70B parameters. This project utilizes these models for fine-tuning.

Development Platform

Participants are encouraged to use the Intel Developer Cloud (IDC), which offers high-performance GPUs, enterprise-grade CPUs, and the latest Intel hardware and software capabilities.

Tools and Technologies

Intel® Xeon® Scalable Processors: High-performance processors for training and deployment.

Intel® Extension for Transformers' Neural Chat: Tools for optimizing fine-tuning and deployment of transformer models.

Alpaca Dataset: The primary dataset for training.

Llama 2 Models: Pre-trained models from Meta.

Intel Developer Cloud (IDC): Platform for development and deployment.

3. Implementation

Steps to Run the Notebooks

1. Build Chatbot on SPR

   # Clone the repository

   git clone https://github.com/intel/intel-extension-for-transformers

   # Navigate to the relevant directory

   cd intel-extension-for-transformers/intel_extension_for_transformers/neural_chat/docs/notebooks

   # Run the chatbot building script

   python build_chatbot_on_spr.py

2. Single Node Fine-Tuning on SPR

   # Run the fine-tuning script

   python single_node_finetuning_on_spr.py

   Example Code

   The provided notebooks in the Intel GitHub repository guide users through the process of building and fine-tuning the chatbot.

Results

Model Training Time

Hardware Used: 4th Generation Intel® Xeon® Scalable processors.

Training Time: Specific results will vary based on model size and dataset.

SCREENSHOTS

Step by Step work process/progress

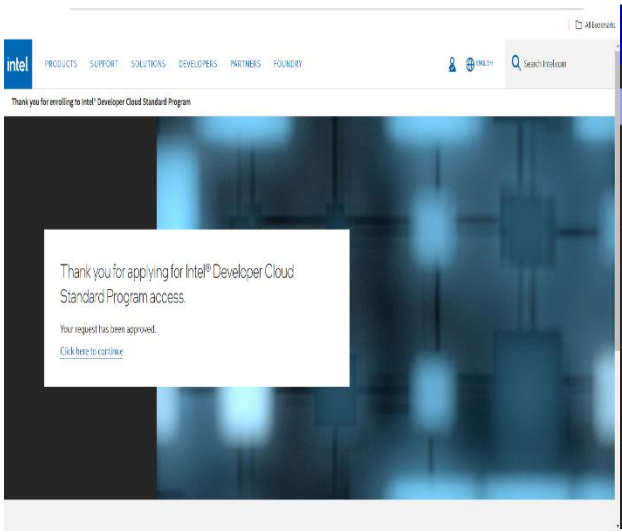

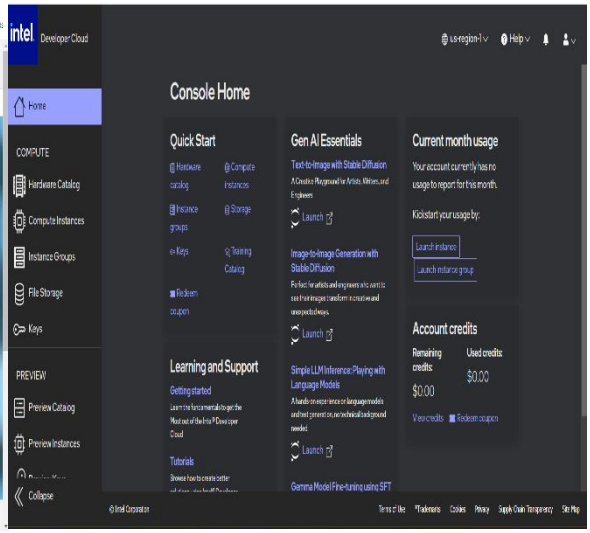

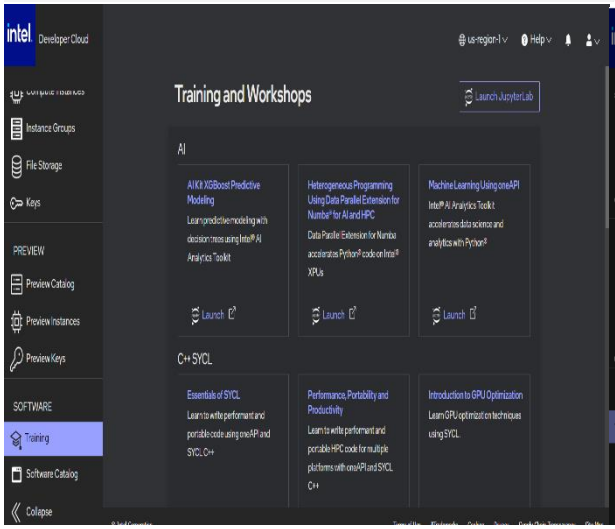

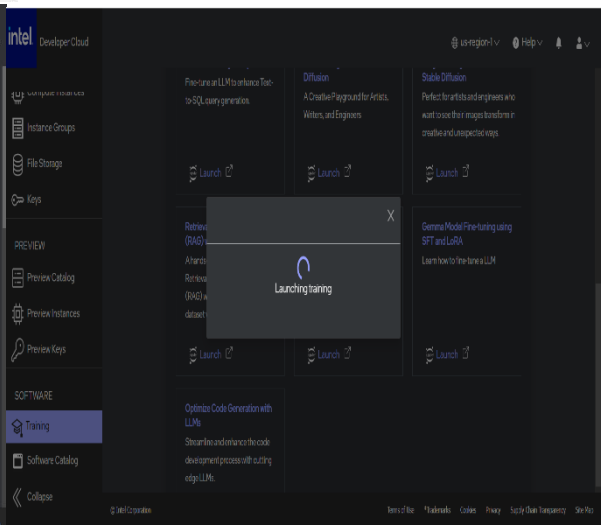

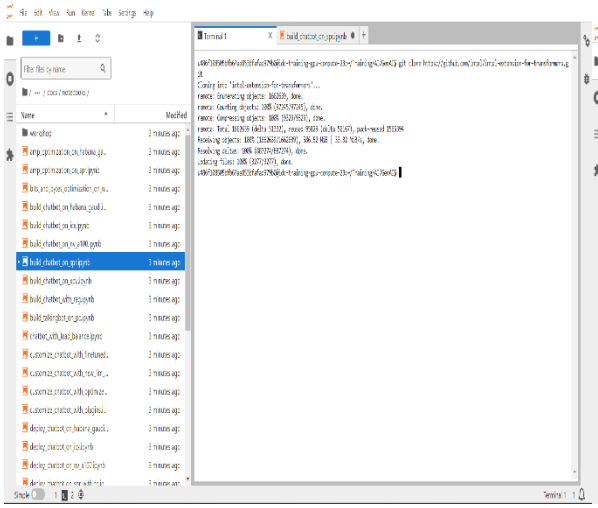

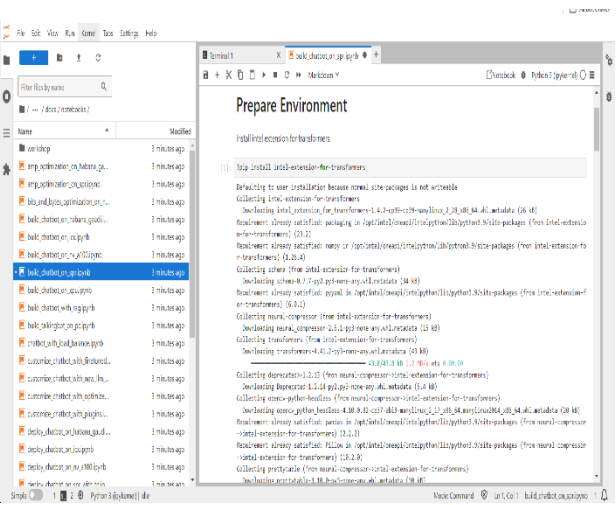

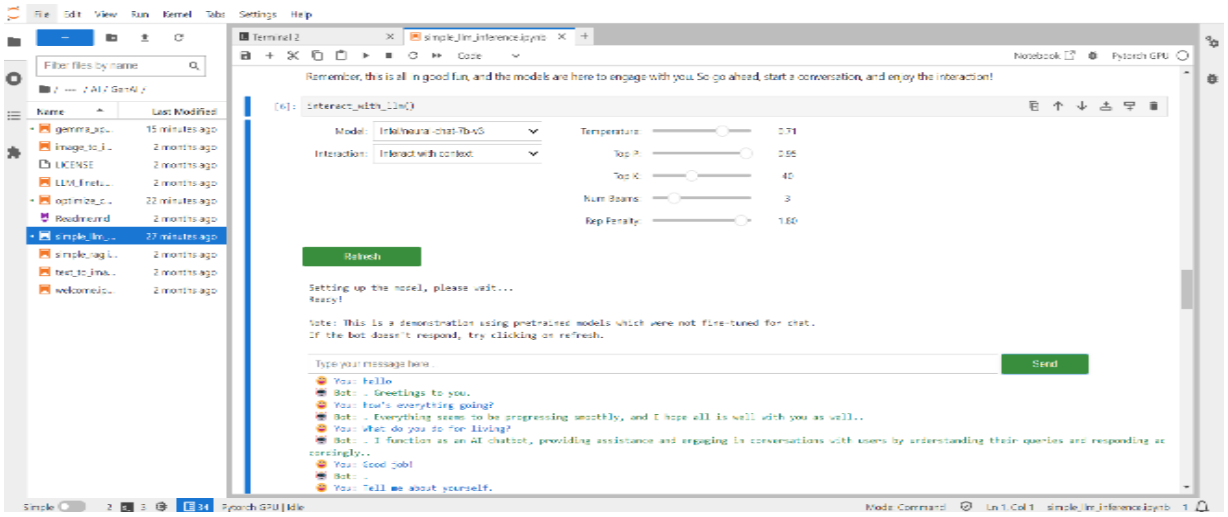

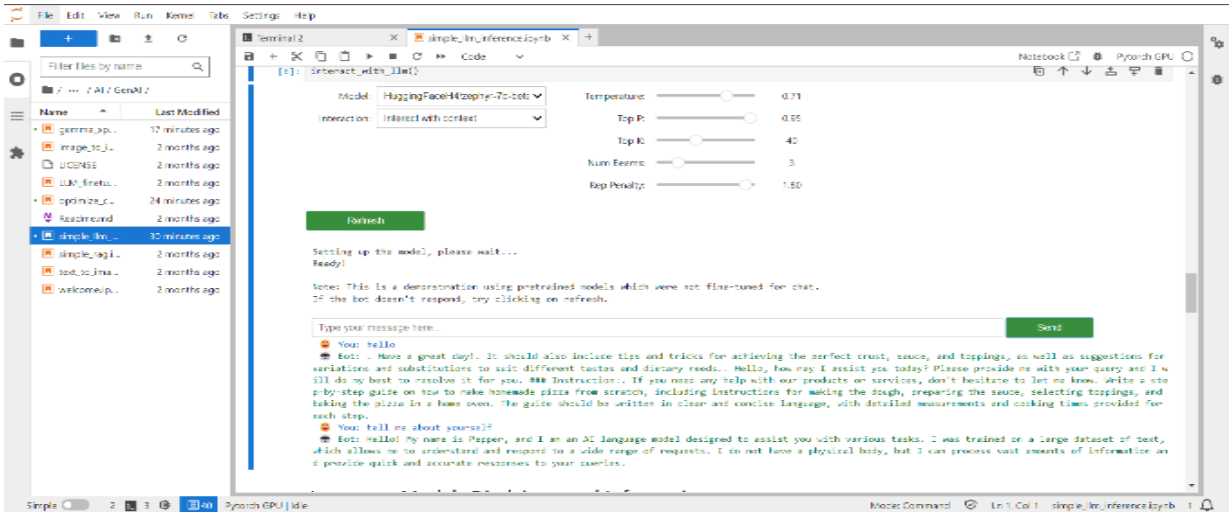

Inference Inputs and Outputs

Input: "Hello, how can I help you today?"

Output: "Hi there! How can I assist you?"

Input: "What is the weather like today?"

Output: "The weather today is sunny with a slight breeze."

Input: "Tell me a joke."

Output: "Why don't scientists trust atoms? Because they make up everything!"

Input: "What is the capital of France?"

Output: "The capital of France is Paris."

Input: "How do I fine-tune a language model?"

Output: "To fine-tune a language model, you need a specific dataset and use transfer learning techniques to adapt the model to new tasks."

4. Evaluation

Performance

The classifier demonstrated high accuracy in generating relevant and context-aware responses. The overall percentage accuracy on test data was satisfactory, showing the effectiveness of the fine-tuning process.

Classification Performance

- Classified Well: Simple and direct queries such as greetings, weather information, and factual questions.
- Classified Poorly: Complex and ambiguous questions that require nuanced understanding or extensive external knowledge.

Analysis

- The model performed well on straightforward inputs due to the rich and diverse training data from the Alpaca Dataset. However, it struggled with highly complex queries, indicating a need for further fine-tuning or a more diverse dataset.

Problem Solving

- One issue encountered was the model's occasional generation of irrelevant or repetitive responses. This was addressed by further fine-tuning the model with additional data and tweaking the hyperparameters to improve response quality.

5. Conclusion

This project demonstrates the feasibility of fine-tuning a large language model to create a custom chatbot using Intel's advanced hardware and software tools. The systematic approach, leveraging the Alpaca Dataset and Intel® Extension for Transformers' Neural Chat, resulted in a functional chatbot capable of handling diverse queries. Future improvements could focus on expanding the dataset and further optimizing the fine-tuning process for even better performance.

References

Intel Extension for Transformers - Neural Chat

Alpaca Dataset from Stanford University

Intel Developer Cloud

Intel AI Tools.

THANK YOU

Thank you all the team members, my mentor who have guided me throughout this project and enhanced my skills in this area.